

System for Emotional Response and Analysis (SERS): A Classroom Emotion Intelligence Platform

1st K. Deepa Shree
Assistant Professor
*Department of Computer Science
& Engineering*
*Dayananda Sagar Academy of
Technology and Management*
*Bengaluru, Karnataka, India -
560082*
deepashree-cse@dsatm.edu.in

4th Inarah Fatima Shariff
*Computer Science and
Engineering*
*Dayananda Sagar Academy of
Technology and Management*
Bangalore, India
Inarahshariff@gmail.com

2nd Honmayee SP
*Computer Science and
Engineering*
*Dayananda Sagar Academy of
Technology and Management*
Bangalore, India
honmayeesp@gmail.com

5th Deepthi R
*Computer Science and
Engineering*
*Dayananda Sagar Academy of
Technology and Management*
Bangalore, India
deepthir1105@gmail.com

3rd Gnanitha P
*Computer Science and
Engineering*
*Dayananda Sagar Academy of
Technology and Management*
Bangalore, India
gnanitha19p@gmail.com

Abstract—

In modern classroom environments, it is often difficult for teachers to continuously monitor how students respond to a lesson while simultaneously managing instruction and interaction. Subtle indicators such as confusion, disengagement, or anxiety may not be immediately visible, especially in larger classes.

This paper presents the System for Emotional Response and Analysis (SERS), a classroom-support platform designed to provide real-time insights into student emotional trends. The system integrates computer vision techniques with a full-stack architecture, where facial features are detected using OpenCV and analyzed through the DeepFace framework. A Django-based backend processes and organizes the data, while a React-based dashboard presents simplified, teacher-friendly visualizations.

To address privacy concerns, the system avoids storing raw classroom images and instead retains only processed emotion labels, timestamps, and confidence scores. The primary objective of SERS is not to replace teacher judgment, but to assist in identifying patterns such as sustained confusion or reduced engagement during a session. By presenting aggregated trends rather than isolated predictions, the platform aims to support timely and informed teaching interventions.

The paper outlines the system design, methodology, implementation approach, and key considerations related to usability, ethics, and data privacy.

I. INTRODUCTION

Learning in a classroom is influenced not only by what is taught, but also by how students respond during the teaching process. In many cases, a student may remain silent while experiencing confusion, loss of focus, or hesitation to ask questions. From a teacher's perspective, especially in larger classrooms, these subtle differences are often difficult to observe in real time.

Traditional indicators such as tests, assignments, and attendance records provide useful information about student

performance. However, they usually reflect outcomes after the learning process has already taken place. This delay limits the ability of teachers to respond immediately when students begin to struggle with a concept.

Recent developments in affective computing and facial-expression analysis have opened up possibilities for observing visible emotional cues during classroom sessions. When used carefully, such technologies can help identify patterns related to engagement, confusion, or discomfort. However, interpreting these signals in an educational setting requires caution, as emotional expressions may vary across individuals and situations.

The System for Emotional Response and Analysis (SERS) is proposed as a classroom-support platform that aims to assist teachers by providing simplified insights into classroom emotional trends. The system integrates computer vision techniques with a full-stack architecture, where facial cues are analyzed and converted into aggregated patterns rather than isolated predictions.

One of the key motivations behind this work is the need for timely intervention. For instance, if a significant portion of the class appears confused during a particular topic, the teacher can adjust the pace, provide additional explanations, or introduce interactive elements. In this way, the system acts as a supportive layer rather than an automated decision-making tool.

Another important consideration is usability. Raw model outputs are often difficult to interpret during a live class. Therefore, SERS focuses on presenting information in a simplified and actionable form, such as trends and summaries, which can be quickly understood without interrupting the teaching flow.

The scope of the current work is limited to observable facial cues captured in a controlled classroom environment. It does not attempt to evaluate academic ability, behavioral patterns, or psychological conditions. Instead, the system is

designed as an assistive tool, where all outputs remain subject to human interpretation and ethical considerations.

Overall, this paper presents the design and development of SERS as a practical integration of emotion recognition with a complete full-stack workflow, aimed at improving classroom awareness while maintaining privacy and responsible usage.

II. RELATED WORK

Research in affective computing has provided the foundation for systems that attempt to recognize and interpret human emotions using computational methods. Early work in this area focused on defining emotional categories and understanding how they can be represented through observable signals such as facial expressions. With the introduction of publicly available datasets like FER-2013 and AffectNet, it became easier to train and evaluate deep learning models for facial-expression recognition.

In recent years, advances in deep learning have significantly improved the performance of visual recognition systems. Convolutional neural networks, in particular, have been widely used for extracting features from facial images and classifying emotional states. Frameworks such as DeepFace offer ready-to-use tools for facial analysis, making them suitable for rapid prototyping and experimental systems. Similarly, libraries like OpenCV continue to play an important role in tasks such as face detection and preprocessing, although their effectiveness can vary depending on real-world conditions such as lighting and camera angles.

While these technologies have shown strong results in controlled environments, applying them directly to classroom settings introduces additional challenges. Factors such as occlusion, varying seating positions, and differences in individual expression can affect the reliability of predictions. Moreover, many existing approaches focus primarily on model accuracy, without addressing how the output will be interpreted or used by teachers during an actual class session.

Educational research has also explored the use of facial cues and attention indicators to understand student engagement. These studies highlight the potential of using visual signals to support learning analytics. However, practical deployment requires more than accurate detection. The information must be presented in a way that is easy to understand, ethically responsible, and useful for decision-making within the classroom.

Another limitation observed in prior work is the lack of integration between emotion-recognition models and complete application systems. In many cases, the focus remains on the prediction stage, without extending the solution to include data storage, visualization, or actionable insights. As a result, there is a gap between technical capability and real-world usability.

Table I. Comparative summary of related work and identified research gap

Reference Area	Primary Focus	Limitation Observed	How SERS Addresses the Gap
Affective computing [3], [6]	Theoretical foundation for recognizing and interpreting emotional states.	Does not directly define a deployable classroom workflow.	SERS converts affective concepts into teacher-facing alerts and reports.
FER datasets [4], [5]	Model training and benchmarking for facial-expression recognition.	Dataset accuracy alone does not ensure classroom usability.	SERS combines model output with aggregation, trend logic, and human review.
DeepFace framework [7]	Facial attribute and emotion analysis using pre-trained models.	Acts mainly as an AI component, not a full educational platform.	SERS integrates DeepFace into a Django-based full-stack system.
OpenCV face detection [2], [8]	Efficient face localization in images and video frames.	Detection quality depends on camera angle, lighting, and occlusion.	SERS adds preprocessing, confidence filtering, and trend-level interpretation.
Classroom engagement studies [9], [10]	Detection of student engagement and attention cues.	Limited focus on notification, privacy, and teacher coaching.	SERS adds coaching suggestions, privacy-aware logging, and parent communication logic.

The proposed SERS platform addresses this gap by combining emotion recognition with a full-stack implementation that includes backend processing, structured data storage, dashboard visualization, and rule-based interpretation. Instead of presenting raw predictions, the system emphasizes aggregated trends and simplified outputs, which are more suitable for classroom use. This approach aims to bridge the gap between theoretical models and practical application while maintaining attention to privacy and ethical considerations.

III. SYSTEM ARCHITECTURE AND METHODOLOGY

The proposed system follows a structured workflow in which classroom video input is processed and converted into meaningful insights for the teacher. Instead of focusing only on model predictions, the design emphasizes how the data moves through different stages and how it is ultimately presented in a usable form.

At a high level, the process begins with video capture, followed by face detection, preprocessing, emotion classification, data storage, and visualization. These stages are connected through a backend layer that manages communication between components and ensures that the system operates in a coordinated manner.

The architecture is intentionally designed to be modular. Each component, such as face detection or emotion classification, can be modified or replaced without affecting the rest of the system. This makes the design more flexible for experimentation and future improvements, especially in a student project setting.




Fig. 1. Proposed SERS system architecture with Django as the central backend layer.

A. Data Acquisition

The system uses a classroom camera to capture video from a fixed position. The setup assumes moderate lighting and a viewing angle that allows most students to be visible without requiring close-up capture. To reduce unnecessary data processing, the system samples frames at regular intervals instead of analyzing every frame in the video stream. This approach balances computational efficiency with the need to capture meaningful trends over time.

B. Face Detection and Preprocessing

Captured frames are processed to identify faces using OpenCV-based detection techniques. Once detected, facial regions are extracted and normalized before being passed to the emotion recognition stage. This step helps maintain consistency despite variations in distance, lighting, and positioning within the classroom.

However, real-world conditions can still affect detection quality. Situations such as partial occlusion, side-facing angles, or poor lighting may reduce reliability. For this reason, the system considers confidence levels and avoids treating uncertain detections as strong indicators.

C. Emotion Classification

The preprocessed facial data is analyzed using a DeepFace-based model to estimate emotional states. The model typically produces multiple emotion categories, but for classroom use, these outputs are grouped into simpler and more interpretable labels such as neutral, happy, confused, or anxious.

This simplification is important because classroom decisions are rarely based on fine-grained distinctions. Instead, broader patterns—such as increasing confusion or sustained low engagement—are more useful for guiding teaching adjustments.

D. Data Logging

To address privacy concerns, the system does not store raw images or video recordings. Instead, it records only essential information such as timestamps, session identifiers, predicted emotion labels, and confidence values. These records are stored in a structured format using database models managed by the backend.

This approach follows the principle of data minimization, ensuring that the system retains only the information necessary for analysis while reducing potential risks associated with sensitive data storage.

E. Dashboard and Alerts

The processed data is presented through a web-based dashboard that displays summarized classroom trends. Rather than showing raw model outputs, the interface highlights key patterns such as dominant emotional states and changes over time.

The design focuses on clarity and quick interpretation. For example, charts and summary indicators allow teachers to understand the overall classroom response without interrupting the flow of teaching. This makes the system more practical for real-time use as well as post-session review.

F. Teacher Coaching and Parent Notification

Based on observed patterns, the system can generate simple suggestions to assist teachers. For instance, an increase in confusion may indicate the need to revisit a topic or provide additional examples. These suggestions are not intended to replace teacher decisions but to offer supportive guidance.

Any form of notification, especially those involving parents, is designed to remain controlled and subject to institutional policies. This ensures that sensitive information is handled responsibly and reviewed before being shared.

Overall, the methodology emphasizes not only accurate detection but also meaningful interpretation and responsible usage. By combining multiple components into a single workflow, the system moves beyond isolated prediction and focuses on practical classroom support.

10-Step Workflow




Fig. 2. Ten-step workflow followed by the SERS prototype from data acquisition to reporting.

G. Django-Based Backend Architecture

The backend of the system is implemented using Django and Django REST Framework, which together handle data processing, storage, and communication between different components. Instead of functioning as a simple data receiver, the backend acts as a central layer that manages how information flows from the detection pipeline to the user interface.

The system uses structured database models such as Student, ClassSession, EmotionRecord, AlertRecord, and ReportSummary to organize processed emotion data. These models make it easier to maintain consistency, retrieve session-based information, and generate summaries when required.

To enable interaction with the frontend, the backend exposes RESTful API endpoints for different functionalities. For example, recent emotion records, aggregated trends, alerts, and reports are accessed through dedicated endpoints. The React-based dashboard retrieves this data and presents it in the form of charts, summaries, and simple recommendations for the teacher.

This API-driven approach follows standard REST principles, allowing the frontend and backend to remain loosely coupled. As a result, changes in one layer can be managed without significantly affecting the other.

For computationally intensive tasks such as emotion inference, the system can optionally use asynchronous processing through tools like Celery and Redis. This prevents delays in user-facing operations and ensures that dashboard interactions remain responsive.

The same architectural design also supports scalability. During deployment, components such as the web server, background worker, and database can be separated, making the system more adaptable to increased usage or future extensions.

IV. TECHNOLOGY STACK

The selection of technologies for the SERS platform is guided by the need to balance simplicity, performance, and practical implementation within a classroom environment. Rather than focusing on highly complex tools, the system uses widely adopted frameworks that are easier to understand, integrate, and extend.

For face detection and preprocessing, OpenCV is used due to its reliability and ease of integration with Python-based applications. It provides sufficient performance for detecting faces in real-time classroom scenarios, particularly in controlled environments.

Emotion analysis is performed using the DeepFace framework, which offers pre-trained models for facial emotion recognition. This allows the system to estimate emotional states without requiring extensive model training, making it suitable for a prototype-level implementation.

The backend is developed using Django along with Django REST Framework, which together handle data processing, API communication, and database management. This combination provides a structured and scalable approach for building web-based systems while simplifying tasks such as authentication and data handling.

On the frontend, React.js is used to build an interactive dashboard that presents classroom insights in a clear and accessible manner. Visualization libraries such as Chart.js are integrated to display emotion distributions and trends over time.

For data storage, the system relies on relational databases such as PostgreSQL or SQLite, managed through Django's ORM. This ensures that emotion records and session data are stored in a structured format that supports efficient querying and reporting.

Additional components such as Firebase Cloud Messaging can be used for notifications, while report generation is handled through PDF-based summaries. These features enhance the usability of the system by enabling communication and documentation of classroom observations.

Another important aspect of the chosen technology stack is its accessibility for development and testing. Since the system is intended as a prototype within an academic environment, the use of widely available frameworks such as Django, React, and OpenCV reduces the learning curve and allows for faster implementation. In addition, these technologies are well-supported by documentation and community resources, which makes debugging and further enhancement more manageable. This practical choice of tools ensures that the system remains maintainable and can be easily extended in future work without requiring significant changes to the core architecture.

Overall, the technology stack reflects a full-stack approach in which each layer—data processing, backend logic, storage, and visualization—works together to deliver a coherent and practical classroom-support system.




Fig. 3. Layered technology stack of SERS, highlighting Django and Django REST Framework as the core backend.

V. RESULTS AND OBSERVATIONS

The prototype was evaluated using recorded classroom-style video and controlled sample inputs. These results should be interpreted as preliminary observations, since real classroom deployment would require additional approvals and more diverse testing conditions.

Under favorable conditions, such as adequate lighting and a clear front-facing camera view, the system was able to detect faces consistently and generate stable emotion predictions. However, the performance varied when students were positioned farther from the camera or when faces were partially occluded. This indicates that practical deployment will depend significantly on environmental factors such as camera placement and lighting.

One of the key observations was the usefulness of aggregated data over time. Individual frame-level predictions were often inconsistent, but when grouped across multiple frames and students, clearer patterns began to emerge. For example, a gradual increase in confusion-related labels during a topic provided a more meaningful signal than isolated predictions.

The dashboard played an important role in improving usability. By converting multiple predictions into simple visual summaries, it allowed the overall classroom state to be understood quickly. Instead of focusing on individual students, the system highlights broader trends such as dominant emotional states or changes over time.




Fig. 4. Sample emotion distribution chart for classroom-level monitoring.

From a design perspective, the decision to avoid storing raw images proved beneficial for privacy. Since only processed emotion labels and timestamps are stored, the system reduces the risk associated with handling sensitive visual data. At the same time, it retains enough information to support meaningful analysis.




Fig. 5. Sample time-window trend graph showing how alerts can support teacher intervention.

Another important observation is that emotion recognition should not be treated as a definitive measure of student understanding. Facial expressions can be influenced by many external factors, and therefore the results are best used as supportive indicators rather than final conclusions. This reinforces the role of the system as an assistive tool rather than a decision-making system.

In addition to this, the reliability of the system improves when observations are considered over a period of time rather than at a single moment. Temporary expressions or distractions may not accurately reflect a student's engagement with the lesson. By focusing on patterns that persist across multiple frames and multiple students, the system is able to provide more stable and meaningful insights. This time-based approach reduces the chances of misinterpretation and makes the output more relevant for classroom decision-making.

Table II. Evaluation focus used for prototype assessment

Evaluation Aspect	What is Observed	Reason for Inclusion
Face detection reliability	Number of visible faces detected under normal lighting and seating conditions.	Determines whether enough classroom input is available for trend analysis.
Emotion trend consistency	Whether repeated emotion labels form meaningful time-window patterns.	Prevents the system from overreacting to a single frame or temporary expression.
Dashboard readability	Clarity of emotion cards, charts, alerts, and teacher suggestions.	Ensures that teachers can use the output during or after class.
Privacy protection	Absence of stored raw images and use of pseudonymised records.	Reduces risk when handling sensitive student-related data.

Overall, the evaluation highlights that while the system performs reasonably well under controlled conditions, its effectiveness in real environments will depend on careful implementation, user awareness, and continuous refinement.

VI. SECURITY, PRIVACY, AND ETHICAL CONSIDERATIONS

Since the proposed system deals with student-related data, security and privacy are critical aspects of the design. The system is intended to operate within institutional guidelines, where proper authorization and transparency are maintained throughout its usage. Students should be informed about how the system works, what data is being collected, and how it will be used.

To minimize risk, the system avoids storing raw images or continuous video recordings. Instead, it retains only processed information such as emotion labels, timestamps, and confidence values. This approach follows the principle of data minimization, ensuring that only the necessary data is stored for analysis while reducing the chances of misuse.

From a security perspective, the backend can implement standard protection mechanisms such as authentication, role-based access control, and secure communication protocols. These measures help ensure that only authorized users, such as teachers or administrators, can access the system's data and functionality.

Ethical considerations are equally important, as emotional data can be sensitive and subject to misinterpretation. The system is not designed to label or judge students, but rather to assist teachers in identifying general classroom patterns. Any insights generated by the system should be reviewed by a human before taking action.

Furthermore, the system should not be used for surveillance or disciplinary purposes. Its primary role is to support learning by highlighting potential areas where students may need additional attention. Responsible use, combined with institutional policies and human oversight, is essential to ensure that the system benefits students without compromising their privacy or well-being.

VII. LIMITATIONS AND FUTURE WORK

Although the proposed system demonstrates promising results, it has several limitations that must be acknowledged. The accuracy of emotion recognition is highly dependent on external factors such as lighting conditions, camera angle, and the visibility of faces. In real classroom environments, issues such as occlusion, movement, and varying seating arrangements can affect detection reliability. In addition, facial expressions alone may not always accurately represent a student's level of understanding, as emotions can be influenced by unrelated personal or environmental factors.

Another limitation is that the current system operates within a controlled setup and has not been extensively tested in diverse classroom scenarios. Differences in student behavior, cultural expression, and classroom dynamics may impact the consistency of results. Therefore, the system should be treated as an experimental prototype rather than a fully validated solution.

Future work will focus on improving the robustness of the system by incorporating more advanced detection techniques and optimizing performance under real-world conditions. The integration of additional data sources, such as voice tone or interaction patterns, could provide a more comprehensive understanding of classroom engagement. However, such extensions must be carefully balanced with privacy considerations.

Further improvements can also include the use of faster processing techniques and hardware acceleration to support real-time performance in larger classrooms. In addition, incorporating feedback from teachers can help refine the system's suggestions and make it more aligned with practical teaching needs.

Overall, future development should aim to enhance accuracy, scalability, and usability while maintaining strong ethical and privacy safeguards.

VIII. CONCLUSION

This paper presented the design and development of the System for Emotional Response and Analysis (SERS), a classroom-support platform that integrates emotion recognition with a full-stack application framework. The system combines computer vision techniques with backend processing and interactive visualization to provide insights into classroom emotional trends.

The primary objective of the proposed approach is to assist teachers by highlighting patterns such as confusion or reduced engagement during a session. Rather than focusing on individual predictions, the system emphasizes aggregated trends that are more meaningful in a classroom context. This makes the output easier to interpret and more relevant for timely teaching interventions.

From our observations during testing we noticed an important aspect of the design is its focus on privacy and responsible usage. By avoiding the storage of raw images and relying on processed data, the system attempts to reduce risks associated with handling sensitive information. At the same time, the role of the system is clearly defined as

assistive, ensuring that all decisions remain under human supervision.

Although the current implementation is limited to a prototype, the results demonstrate the potential of combining emotion recognition with full-stack development for practical educational applications. With further testing and ethical consideration, the system can be extended to support more reliable and scalable classroom environments.

Overall, SERS highlights how technology can be used to improvise classroom awareness while maintaining a balance between functionality, usability, and privacy.

IX. ACKNOWLEDGMENT

The authors would like to express their sincere gratitude to Prof. K. Deepa Shree for her continuous guidance, valuable suggestions, and support throughout the development of this project. Her insights played an important role in shaping the direction of this work.

We also thank the Department of Computer Science and Engineering at Dayananda Sagar Academy of Technology and Management for providing the necessary resources and environment to carry out this project successfully.

Finally, we would like to acknowledge all those who contributed directly or indirectly to the completion of this work.

X. REFERENCES

- [1] S. Li and W. Deng, "Deep facial expression recognition: A survey," *IEEE Transactions on Affective Computing*, vol. 13, no. 3, pp. 1195–1215, 2022.
- [2] P. Viola and M. Jones, "Rapid object detection using a boosted cascade of simple features," in *Proc. IEEE CVPR*, 2001, pp. 511–518.
- [3] R. W. Picard, *Affective Computing*. Cambridge, MA, USA: MIT Press, 1997.
- [4] I. J. Goodfellow et al., "Challenges in representation learning: A report on three machine learning contests," *Neural Networks*, vol. 64, pp. 59–63, 2015.
- [5] A. Mollahosseini, B. Hasani, and M. H. Mahoor, "AffectNet: A database for facial expression, valence, and arousal computing in the wild," *IEEE Transactions on Affective Computing*, vol. 10, no. 1, pp. 18–31, 2019.
- [6] P. Ekman, "An argument for basic emotions," *Cognition and Emotion*, vol. 6, nos. 3–4, pp. 169–200, 1992.
- [7] S. I. Serengil and A. Ozpinar, "HyperExtended LightFace: A facial attribute analysis framework," in *Proc. Int. Conf. Engineering and Emerging Technologies*, 2021, pp. 1–4.
- [8] G. Bradski, "The OpenCV Library," *Dr. Dobbs's Journal of Software Tools*, 2000.
- [9] J. Whitehill et al., "The faces of engagement: Automatic recognition of student engagement from facial expressions," *IEEE Transactions on Affective Computing*, vol. 5, no. 1, pp. 86–98, 2014.
- [10] T. Liao, "Gaze-based attention estimation for e-learning," *Computers & Education*, vol. 157, Article 103989, 2020.
- [11] M. Pantic and L. J. M. Rothkrantz, "Automatic analysis of facial expressions: The state of the art," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 22, no. 12, pp. 1424–1445, 2000.
- [12] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet classification with deep convolutional neural networks," *Communications of the ACM*, vol. 60, no. 6, pp. 84–90, 2017.
- [13] J. Deng et al., "RetinaFace: Single-shot multi-level face localisation in the wild," in *Proc. IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020, pp. 5203–5212.
- [14] Django Software Foundation, "Django documentation," 2024.
- [15] R. T. Fielding, "Architectural styles and the design of network-based software architectures," Ph.D. dissertation, University of California, Irvine, 2000.
- [16] OWASP Foundation, "OWASP Top 10: The ten most critical web application security risks," 2021.
- [17] European Parliament and Council of the European Union, "General Data Protection Regulation (GDPR)," 2016.
- [18] S. Slade and P. Prinsloo, "Learning analytics: Ethical issues and dilemmas," *American Behavioral Scientist*, vol. 57, no. 10, pp. 1510–1529, 2013.
- [19] M. Soleymani, M. Pantic, and T. Pun, "Multimodal emotion recognition in response to videos," *IEEE Transactions on Affective Computing*, vol. 3, no. 2, pp. 211–223, 2012.
- [20] M. F. Valstar et al., "FERA 2015 – Facial Expression Recognition and Analysis Challenge," in *Proc. IEEE Int. Conf. Automatic Face and Gesture Recognition*, 2015, pp. 1–8.